


Start coding or [generate](#) with AI.

```
from google.colab import drive
drive.mount('/content/drive')
```

 Mounted at /content/drive

✓ Image Classification with ANN Using Fashion MNIST Dataset

Project Type: EDA and ANN

Objective

To build an Artificial Neural Network that classifies images of clothing items (like t-shirts, trousers, and shoes) into predefined categories.

Double-click (or enter) to edit

Dataset Info:

The Fashion MNIST dataset is a set of 28x28 grayscale images of 10 fashion categories:

- T-shirt/top
- Trouser
- Pullover
- Dress
- Coat
- Sandal
- Shirt
- Sneaker
- Bag
- Ankle boot


Importing libraries

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from sklearn.metrics import classification_report, confusion_matrix
```

Start coding or [generate](#) with AI.

Load and Preprocess the Data

```
fashion_mnist_dataset = tf.keras.datasets.fashion_mnist
(X_train, y_train), (X_test, y_test) = fashion_mnist_dataset.load_data()
```

 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
 29515/29515 ————— 0s 0us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
 26421880/26421880 ————— 0s 0us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
 5148/5148 ————— 0s 1us/step
 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
 4422102/4422102 ————— 0s 0us/step

Normalize the pixel values to be between 0 and 1

```
X_train = X_train / 255.0
X_test = X_test / 255.0
```

✓ Exploratory Data Analysis (EDA)

```
# Plot the first 10 images in the training dataset
plt.figure(figsize=(10,10))
for i in range(10):
    plt.subplot(2, 5, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(X_train[i], cmap=plt.cm.binary)
    plt.xlabel(f'Label: {y_train[i]}')
plt.show()
```



Label: 9




Label: 0



Label: 0




Label: 3




Label: 0




Label: 2




Label: 7



Label: 2



Label: 5



Label: 5

Each image corresponds to a clothing category with a label (0–9), which we will map as follows:

- 0 = T-shirt/top
- 1 = Trouser
- 2 = Pullover
- 3 = Dress
- 4 = Coat
- 5 = Sandal
- 6 = Shirt
- 7 = Sneaker
- 8 = Bag
- 9 = Ankle boot

Distribution of Classes

```
# Distribution of the labels in the training dataset
unique, counts = np.unique(y_train, return_counts=True)
class_distribution = dict(zip(unique, counts))
```

```
print("Class Distribution: ", class_distribution)
```

```
↗ Class Distribution: {0: 6000, 1: 6000, 2: 6000, 3: 6000, 4: 6000, 5: 6000, 6: 6000, 7: 6000, 8: 6000, 9: 6000}
```

Build the Artificial Neural Network (ANN)

```
# Initialize the model
model = Sequential()

# Flatten the 28x28 images into a 1D array (28*28 = 784)
model.add(Flatten(input_shape=(28, 28)))

# Add hidden layers with ReLU activation
model.add(Dense(units=128, activation='relu'))
model.add(Dense(units=64, activation='relu'))

# Output layer with 10 units for the 10 classes, using softmax activation
model.add(Dense(units=10, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Model summary
model.summary()
```

```
↗ /usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim`
  super().__init__(**kwargs)
Model: "sequential"
```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100,480
dense_1 (Dense)	(None, 64)	8,256
dense_2 (Dense)	(None, 10)	650

Total params: 109,386 (427.29 KB)
Trainable params: 109,386 (427.29 KB)
Non-trainable params: 0 (0.00 B)

Train the Model

```
# Train the model
history = model.fit(X_train, y_train, validation_split=0.2, epochs=10,
                    batch_size=32)
```

```
↗ Epoch 1/10
1500/1500 ————— 10s 5ms/step - accuracy: 0.7638 - loss: 0.6819 - val_accuracy: 0.8601 - val_loss: 0.3887
Epoch 2/10
1500/1500 ————— 10s 5ms/step - accuracy: 0.8583 - loss: 0.3886 - val_accuracy: 0.8586 - val_loss: 0.3883
Epoch 3/10
1500/1500 ————— 8s 4ms/step - accuracy: 0.8731 - loss: 0.3452 - val_accuracy: 0.8634 - val_loss: 0.3802
Epoch 4/10
1500/1500 ————— 10s 4ms/step - accuracy: 0.8805 - loss: 0.3209 - val_accuracy: 0.8674 - val_loss: 0.3590
Epoch 5/10
1500/1500 ————— 8s 6ms/step - accuracy: 0.8893 - loss: 0.2979 - val_accuracy: 0.8802 - val_loss: 0.3313
Epoch 6/10
1500/1500 ————— 6s 4ms/step - accuracy: 0.8964 - loss: 0.2809 - val_accuracy: 0.8731 - val_loss: 0.3544
Epoch 7/10
1500/1500 ————— 6s 4ms/step - accuracy: 0.9036 - loss: 0.2597 - val_accuracy: 0.8863 - val_loss: 0.3227
Epoch 8/10
1500/1500 ————— 7s 5ms/step - accuracy: 0.9021 - loss: 0.2565 - val_accuracy: 0.8880 - val_loss: 0.3116
Epoch 9/10
1500/1500 ————— 10s 5ms/step - accuracy: 0.9090 - loss: 0.2447 - val_accuracy: 0.8751 - val_loss: 0.3753
Epoch 10/10
1500/1500 ————— 6s 4ms/step - accuracy: 0.9106 - loss: 0.2374 - val_accuracy: 0.8870 - val_loss: 0.3192
```


Evaluate the Model

```
# Plot accuracy and loss over epochs
plt.figure(figsize=(12, 4))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='train accuracy')
plt.plot(history.history['val_accuracy'], label='val accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='train loss')
plt.plot(history.history['val_loss'], label='val loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



Performance on Test Data

```
# Evaluate the model on the test data
test_loss, test_accuracy = model.evaluate(X_test, y_test)
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")
```

```
313/313 ————— 1s 3ms/step - accuracy: 0.8810 - loss: 0.3402
Test Accuracy: 87.88%
```

Confusion Matrix and Classification Report

```
# Predict the labels for the test set
y_pred = np.argmax(model.predict(X_test), axis=1)

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print(conf_matrix)

# Classification report
print(classification_report(y_test, y_pred, target_names=[
    'T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
    'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']))
```

```
313/313 ————— 2s 6ms/step
[[867  0 15 22  4  1 85  0  6  0]
```

[6	969	1	20	2	0	1	0	1	0]	
[22	1	786	22	106	0	62	0	1	0]	
[29	3	8	914	23	0	20	0	3	0]	
[1	1	96	44	815	0	42	0	1	0]	
[0	0	0	0	0	967	0	11	1	21]	
[157	0	88	34	70	0	645	0	6	0]		
[0	0	0	0	0	36	0	874	0	90]	
[7	0	2	3	5	4	10	3	966	0]	
[0	0	0	0	0	4	1	10	0	985]]	
										</	